

# The Ecology of Cloud Computing

Alex H. Room

February 9, 2022

## 1 Introduction

In regards to the climate emergency, 'the cloud' is a rather apt metaphor for network computing; we are yet to learn whether it will provide us shade from a burning issue, or stand as an omen for a terrible storm. It has the potential to reduce the climate impact of several industries, but if not managed carefully it could cause our global emissions to skyrocket. This report will survey the current climate impact of scientific computing, the benefits of moving to the cloud from an environmental standpoint, and how to mitigate the impact of computing on our environment.

### 1.1 What is cloud computing?

Cloud computing is simply the ability to use computing resources (like storage or processing space) without having to directly own or manage them. For example, if we wanted to have our own website, we could buy a rack of servers ourselves, plug them into the wall, set them up, and voila: we have a website. But what happens when our website gets popular? What happens if one million people all want to use it at one time? We could keep scaling up, fill our kitchen with server racks, start buying air-conditioning for them, and deal with the rather large electricity bill - an approach that many now-huge websites once had to contend with, but eventually you run out of space and time to look after it all without significant investment.

On the other hand, we could host our website on the cloud by renting servers in a data center somewhere. For example, Amazon's Elastic Compute Cloud (EC2) charges \$0.10 per hour for each computer you require, and only for the hours where the computers are actively running. You can host your website off their datacenter, and Amazon has dedicated people who look after the servers and make sure they're running correctly. As you can see, this is a very attractive alternative to self-hosting, particularly at large scales. However, when you do this, it is easy to forget about the real, physical equipment behind whatever computing you're doing - without care, we can fall into seeing the cloud as an abstract *thing* that processes our data for us, an amorphous blob of circuits beyond concrete concerns. Without mindfulness, the emissions problem can thus easily slip out of our hands.

### 1.2 Some intuition for carbon use

Most estimates of emissions are in the units CO<sub>2</sub>e, or "Carbon dioxide equivalent". This is to provide a standardised unit for emissions from different gases - for example, the climate impact of 1kg of methane is the same as the impact of 25kg of CO<sub>2</sub> [1], so if a process emits 3kg of methane and 4kg of CO<sub>2</sub> it has emissions of  $(3*25)+4 = 79\text{kg CO}_2\text{e}$ . Still, we don't generally have good intuition when it comes to knowing what 1g CO<sub>2</sub>e actually *means*. Here are some benchmarks (the first three taken from Lannelongue et al. [2]):

- In a European car, driving one kilometre emits about 175g CO<sub>2</sub>e.
- Flying from Paris to London emits around 50kg CO<sub>2</sub>e.
- Streaming video for one hour emits around 55g CO<sub>2</sub>e (although this is hard to keep track of; electricity use from data transmission has halved every 2 years on average [3]).

- The entire life-cycle of an iPhone 11 used for 3 years is 72kg CO<sub>2</sub>e (79% of which is from manufacturing). [4]
- One kWh of electricity from the UK grid emits 233g CO<sub>2</sub>e. [5]

### 1.3 A miniature case-study of what can go wrong: NFTs

One recent case of computing and climate that has caught media attention is the non-fungible token (NFT). 'Fungibility' is the ability of an object to have an equivalent copy - for example, gold is fungible (every kilogram of gold is worth exactly the same as every other kilogram of gold) whereas, say, a piece of artwork is not (if I were to make a duplicate of Van Gogh's *Starry Night*, it would not be as valuable as the real painting).

An NFT uses an algorithm to create non-fungible digital data - like a digital certificate of authenticity. It uses a technology called blockchain to make a digital receipt that cannot be copied or faked. A blockchain is essentially a list (of transactions) and a signature - the signature is a line of letters and numbers that encapsulates the contents of the list. If someone tampers with the list, or tries to improperly add something to it, the signature and list won't match and you will be able to tell it has been tampered with. [6] This can be a good thing; artists have found use in it, being able to sell digital art as a limited-stock item (like traditional paintings are) directly, without than having to go through a middleman like an auction house. In fact, in 2020, artists earned \$1.8 million [7] from selling their art as NFTs.

Now, where's the catch? The problem is how this uses blockchain. The generation of this signature takes a lot of computation; it is estimated that one NFT transaction generates about 48kg of CO<sub>2</sub> [8], 14 times the footprint of mailing a physical art print, and 74000 times the footprint of buying something via bank transfer. [9] In fact, if you had turned on a laptop the day the Roman Empire was founded, and used it from then until December 2020, that laptop would've produced less carbon emissions than NFTs did in 2020. [8]

And all this is done for an algorithm that isn't even necessary; an alternate blockchain algorithm called proof-of-stake produces the exact same result with only 1% of the computing power. [10] This case study gives a stunning parable on how computing, what we use it for, and what algorithms we use can have a huge effect on the climate.

## 2 Scientific computing and energy

Scientific computing is, essentially, the process of turning electricity into data. To this end, it can be difficult to reduce the amount of electricity used, as this often involves generating less data. Nonetheless, we can review where current and potential problems lie, and where there is room for improvement; it may be useful to view from an efficiency standpoint, where our "science per unit energy" can be improved.

### 2.1 Processing

The website [www.green-algorithms.org](http://www.green-algorithms.org) [11] is an online tool that lets you see the carbon impact of an algorithm. One may not be surprised to learn that the main independent variable here is runtime. When performing computation, a server's power draw is roughly 4 times what it is when the computer is idle [2]. Efficiency, then, is key here; for example, the bioinformatics algorithm BOLT-LMM has a 73% lower carbon footprint on version 2.3 than it did on version 1. [12]

With the ever-increasing processing power and memory of modern computers, it can be easy to not see efficiency as so mission-critical, especially for small algorithms - shaving one microsecond of runtime off a popular algorithm might not seem like much, but if that algorithm is used a trillion times (which isn't so far-fetched, when you consider algorithms like Fast Fourier Transforms or signal processing tools) you've saved 11.6 days of processing time (and 7kg CO<sub>2</sub>e) total. As your program gets faster, there is also tendency to increase its scope and process larger data. This can create a 'rebound' effect [2] that ends up turning your efficiency improvements into larger energy use. In this regard, restraint should be exercised, alongside

careful consideration of how much data you actually *need* to do good science.

One final point is to consider the programming language used when it comes to efficiency benefits. Heer [13] compares the time taken for various programming languages to estimate  $\pi$  via the Leibniz method. He finds that compiled<sup>1</sup> languages such as Go or C++ are the fastest, taking 6-9.5 seconds, whereas some popular interpreted languages such as R or Python3 take 181-243 seconds. Of course, this is not to say that all code should be written in C++ lest we kill the planet. Not everyone knows every programming language, and there are non-speed-related benefits that put one language ahead of the other (like Python's accessibility and large number of add-on modules, or R being the *de facto* standard for statistics in the life sciences). Merely, we should more heavily consider it when assessing what language to start a project in. Carbon cost may sway the balance when we think about how many saved processing hours are worth the extra development hours<sup>2</sup>.

## 2.2 Data

Once we have processed our data, what do we do with it? Data storage counts for around 10% of data center energy usage, for a total of 20TWh per year [14] - nearly 40x the energy used by all personal computers in the UK in a year. [15] Not to mention the energy used to produce all this storage media - the manufacture of a Dell PowerStore9000T hard drive (one of their largest available) produces approximately 2765kg CO<sub>2</sub>e. [16] A possible and quite successful solution used for data archival is to move unused data to magnetic tapes; with careful discernment of what data needs to be at one's fingertips versus on a library shelf. Using tape can reduce the space and utilities needed for storage sixfold. [17].

This is probably the area where it is easiest to 'de-abstract' our thinking. As humans, we are used to thinking in space, and most of us in our daily lives run up to troubles with our phone, laptop or cloud service running out of space for us just as easily as we run out of space in a bookshelf or spare room. On a corporate scale, we may want to think about the actual usefulness of our archives, and 'declutter' as needed. Do we need to keep data from a 10-year-old study that has since far been superseded? Could we offload older data onto relevant parties (essentially, contacting researchers and saying 'we can give you your data if you like, but we aren't keeping it ourselves')? Maybe we could take a page from libraries and keep usage statistics, shedding datasets that haven't been touched or asked for in years? Perhaps in our heads we may do well to separate 'just in case!' from 'just for when...', and assess whether we are archiving or simply hoarding.

## 2.3 Machine learning

Machine learning is becoming *en vogue* in scientific computing; it is being used for everything from improving early stage drug discovery [18] to solving systems of differential equations. [19] In essence, machine learning consists of two stages. The first is a 'training' stage, where a large dataset is processed by an algorithm in a machine; the second is a 'application' stage where the machine then uses what it has learned to make estimations about new data. For example, if I were training a machine to read handwritten letters (see Deng [20]), the 'training' dataset would be a big set of handwritten letters along with what they represent (so we'd tell the machine 'these 50 handwritten digits are all the letter 'a' etc.), and the machine would use an algorithm to figure out similarities and differences between letters. The 'application' stage would then be giving it a handwritten letter it hasn't seen before and it figuring out what letter it is from the data it gathered during the 'training' stage. This means machine learning can do a lot of busywork that we previously had to employ humans to do, like picking out good candidates for new drugs or typing up handwritten letters; essentially, it could be seen as a way of 'automating intuition'.

This learning stage can take a lot of computation - the MNIST database from Deng [20] has 60,000 training letters. The Enron corpus (Klimt and Yang [see 21]), a dataset of emails used for English language

---

<sup>1</sup>'Compiled' means a language where a *compiler* takes your code and turns it into an executable file, which you then run; 'interpreted' means the program is run directly as it goes along, line-by-line by an *interpreter*.

<sup>2</sup>In particular, code for an experiment that only runs once may not be worth the extra development time of a more difficult language, but code for a long-lasting piece of software may be.

machine learning, contains 700,000 emails. And each of these needs an algorithm to be run on it individually. The total process can take days, and for advanced applications can use huge amounts of electricity. Strubell et al. [22] found that training one model using Google’s state-of-the-art BART machine learning framework takes 79 hours, in the process using 1507kWh<sup>3</sup> (approximately the energy used by a UK household in six months [15]) and 326kg CO<sub>2</sub>e (!). Smaller frameworks, like the T2T framework, used around 27kWh and 5kg CO<sub>2</sub>e (which is still a non-negligible amount).

While machine learning has been an amazing development in several fields, it runs the risk of overuse; some software concepts can almost become fashionable, with not enough consideration put into them. Note that the statistics in the previous paragraph were for *training* a model - that is, we’ve generated over 5kg CO<sub>2</sub>e and *we haven’t even solved any problems yet*. Compared to a non-machine learning approach, it is entirely possible that this lump sum of carbon right at the start will put its footprint far ahead of a non-learning-based software for essentially its entire lifetime - care must be taken that machine learning is the right tool for the job.

## 2.4 Computing as supplement

Declaring the discussion finished where software ends would be myopic. So far we have only spoken about software, as if the scientific process is to magically generate data and push it through an algorithm. Most scientific computing is a two-step cycle, where one alternates between analysing data with software and creating data with laboratory instruments. We set up our experiment, collect data, analyse it, then either collect more data or go home and write a paper, as necessary. Through computing, we can add a third step to this cycle (which absorbs part of the ‘creating data’ step) - simulation.

This could be supplemented with an example.<sup>4</sup> The instruments at STFC’s ISIS facility consume around 1.1MW.[23] There are 38 instruments, so we take a rough average of 20KW per instrument. The Intel Xeon X3430 server, which is the highest-power-drawing core listed on green-algorithms.org [11], draws 23.8W per core when active; for high performance computing, one may want to use 8 or 16 cores - we will go with 16 and take 380W power draw for the whole server. As a further assumption (we are going for a bad case here), we will say this server is in a data centre, so add the overhead of the datacentre (per server), which will make it 532W. This gives us an approximately 1:38 ratio of instrument power use versus computing power use. Now, here we have a fascinating figure indeed; if we run an algorithm for less than 38 hours and as a result need one fewer hour of instrument use, we have saved energy (and thus carbon!).

This statistic can be a useful tool to consider priorities in where this simulation step can be leveraged. For example, if a department is trying to decide which simulation tool to put funding into, one with a 1:40 ratio may be more attractive than one with a 1:5 ratio. Although, one disadvantage is that this statistic doesn’t mean anything on an objective level. For a molecular analysis experiment, 1:40 may be very attractive (as one could feasibly reap those time-saving benefits), whereas for a bioinformatics study the complexity of biological processes like protein interactions would make it a struggle to do a significant simulation in that time.

## 3 Cloud computing

Now we will specifically look at cloud computing; the benefits and drawbacks of moving our data to a dedicated external centre.

---

<sup>3</sup>This calculation accounted for the cost of cooling and maintaining the data center during this time, not just the actual processing.

<sup>4</sup>NB: *just* an example. Various other practical concerns vastly complicate this in real life; this simply illustrates the interaction of scientific equipment and computing.

### 3.1 Data centres

A key factor in considering cloud computing is the idea of 'overhead' - around 40% of the energy use of a datacenter is purely the energy used to cool it so the servers don't overheat, [2] which would not be necessary in a local server (which is small enough to be cooled effectively by similar techniques as those used to cool human beings). However, as datacentres are geared purely towards the processing and storage of data, this can be done more efficiently. In fact, a 2013 case study (Masanet et al. [24]) suggested that by moving common applications used in a company from a local server to the cloud, they could reduce emissions by 87%<sup>5</sup>. Naturally, one large cloud centre can utilise efficiency 'economies of scale' better than dozens of local, small ones. And over time, this has certainly shown; research by Google found that while data centre use increased by 550% between 2010 and 2018, the amount of energy consumed by datacentres only increased by 6%, meaning the proportion of global energy used by datacentres did not increase despite their much greater usage. [25]

If using a commercial data centre, one can find a better datacentre by shopping around for better carbon options; with the choice of any cloud provider in the world, you also have the choice of what country's grid they're plugged into. For example, generating one kWh in Switzerland emits 12g CO<sub>2</sub>e, while in Australia it is 880g. [12]. Another statistic of data centre efficiency is its PUE (Power Usage Effectiveness), which measures efficiency of a datacentre via the following ratio:

$$PUE = \frac{\text{Total facility energy use}}{\text{Server energy use}}.$$

Naturally, the ideal PUE is 1. The average datacentre has a PUE of around 1.67, but efficiency improvements can go much lower; for example, Google's datacentres have an average PUE of 1.1. [26]

The PUE of a data centre can be reduced through various considerations; the first is how to manage heat emissions. Cooling the servers using water cooling rather than air conditioning (in fact, air conditioning has recently come under fire for creating a vicious cycle as the energy from air conditioning heats the planet, which means more air conditioning is required, so on [27]) can be much more efficient. And then, design can be used to further improve this; a common way of cooling data centres is by cooling alternating rows of servers while letting the odd rows heat up, which creates a convection current around the whole room that keeps everything cool. The water cooling itself can also be made more effective using ground (passive) cooling - by piping the water underground before piping it back into the data centre, we can take advantage of the effectively unlimited cold under the earth.

We can also integrate some interesting computing approaches to infrastructure to reduce PUE. Evans and Gao [28] reduced the energy use for cooling at a Google datacentre by 40% through training an AI on centre infrastructure data such as temperatures, power, pump speeds (for water cooling) and then allowing it to control the datacentre's cooling system.

If building one's own datacentre, we can also save energy by using renewable energy sources. The main issue with this is in the case of methods like photovoltaic (solar) power, where their non-constant output could cause problems for a facility which needs to be on 24/7.

And then there's the interplay of the datacentre areas vs the human areas; waste heat from servers can be used to heat up areas containing people rather than separate central heating. One incredibly simple way of doing this is to have areas designed for people on higher storeys, so the heat can rise and do as such.

Finally, when processing data and running algorithms in our datacentre, better management and options for users can improve efficiency. One idea could be a 'queueing' system; if not in a rush to get results as soon as possible, users could be given the option to queue their processing for off-peak times where the system would otherwise be wasting electricity just idling away.

---

<sup>5</sup>Not to mention the further reduction in emissions if employees are then able to work from home; if it's all on the cloud anyway, then why would they need to burn petrol to commute?

### 3.2 Moving costs

When we move our data over to a cloud centre, we must choose how we move it, as well as how we move large data back and forth from user to server in future. Thus, time for an interesting comparison; the carbon use of data transfer versus that of a heavy-goods vehicle. Huge corporations like Google transfer data using 'SneakerNet' - that is, loading hard drives into a van and taking them physically to their destination.<sup>6</sup> As long as travel time is under 57 hours, SneakerNet is more efficient than the internet in terms of transfer speed [29] - for reference, crossing the United States takes 40-50 hours. But what about carbon?

Data transmission emits a flat rate of CO<sub>2</sub>e per GB transferred, regardless of distance of transfer. A vehicle emits a certain amount of CO<sub>2</sub>e per kilometre. From this, we can create a simple equation for how much carbon Sneakernet emits compared to internet transfer (assuming that we are not transferring so much data that we cannot fit it all in one vehicle).

$$SC = E_{vehicle}d - E_{internet}s$$

where  $SC$  is the relative Sneakernet carbon (how much carbon emitted by Sneakernet relative to that emitted by internet transfer),  $E_{vehicle}$  is the emissions of the vehicle (in g CO<sub>2</sub>e/km),  $d$  the distance of transfer (in km),  $E_{internet}$  the emissions of data transmission (in g CO<sub>2</sub>e/GB) and  $s$  the size of the data being transferred (in GB). A 2016 study estimated the energy cost of data transfer as 0.023kWh per gigabyte. [3] In the United Kingdom, this would emit 5g CO<sub>2</sub>e/GB. We will take our vehicle as a Mercedes Sprinter (the van used by Royal Mail), which emits approximately 192.6g CO<sub>2</sub>e/km. [30] We will say our facility is in central London, and our data centre at the Cyxtera Data Centre Campus, a popular data centre campus in Slough - a distance of around 35km. This creates a linear plot from which we can see Sneakernet's carbon emissions break even to the Internet's at a data size of 1348.2GB; that is just one or two regular-sized hard drives! This is even better with electric vehicles; Royal Mail plans to switch to the electric Arrival van, which at most stores 139kWh with a range of 340km [31]. If charged using fossil-fuel generated electricity (the worst-case scenario) this emits 94g CO<sub>2</sub>e/km, meaning we 'break even' on SneakerNet at 658GB, which one could probably fit on a single laptop.

An even better-case scenario is if you are building your data centre on the same campus as your facility. Most data centres have a room known as a *meet-me room*, or MMR, where people can enter and physically connect their personal device to the centre, letting them transfer data quickly in person. In this case, SneakerNet will almost certainly win out over data transfer, as the carbon output of walking across a campus is negligible.

### 3.3 A note on carbon offsets

It is tempting consider balancing out the extra electricity use of our datacentre with carbon offsetting. Carbon offsetting is the act of spending money on things with a negative carbon cost to counteract the emissions of a process. This ranges from practices like giving money to reforestation companies, to slightly more dubious things like investing in carbon sequestration research.

Care must be taken when using carbon offsets. At best, they are an important part in a carbon-reduction portfolio, one which actually puts financial incentive on *keeping* ecosystems rather than destroying them for profit. At worst, they are a way for an organisation to funnel money into reducing their climate guilt. It's important to deeply consider offsetting methods; for example, planting trees sounds like a brilliant idea (a mature tree sequesters 11kg of CO<sub>2</sub> per year [32]), but it takes 20 years before a tree is mature enough to start sequestering any significant amount of carbon [33]. By this point, we may have long passed a point of no return for destructive climate change. And besides, are the partners we're using for offsetting reliable? Are they protecting those trees? If your trees catch fire or are eaten by a deer when they've just sprouted, your offset may go entirely unrealised.

---

<sup>6</sup>A wonderful article comparing SneakerNet to traditional bandwidth here: <https://what-if.xkcd.com/31/>

One could make the analogy that carbon offsetting is the climate equivalent of the payday loan. Sure, your carbon output will balance out (over infinite time), but if you can help it you'd do better to just leave the fossil fuels in the ground in the first place.

### 3.4 Futurism

The future, like offsetting, is not to be relied on; we must think carefully about the environment now, rather than relying on scientists for a *deus ex machina*. That said, there are some promising and interesting developments.

The first is DNA data storage. This is where we artificially create DNA chains to store digital information. Not only is this method several orders of magnitude denser than magnetic tapes, it also can remain stable for hundreds or thousands of years, as well as a greener production life-cycle; its biological nature means it can biodegrade when no longer needed, rather than going to landfill or being incinerated. [34]

Another area of interest is quantum computing - this uses concepts from quantum physics to redo computing in a more efficient way 'from scratch'. The difference that this will make to the environment is completely unimaginable; it is estimated to be capable of quickly working through algorithms that normal computers wouldn't be able to do between now and the end of the universe, which completely changes the lens through which we see computing efficiency. However, we cannot tell if the infrastructure required will balance it out in early days; as it stands, quantum computers must be supercooled to work, which uses a massive amount of cooling energy even on a very small scale.

## 4 Conclusion

The main idea flowing underneath this report is one of care; one of mindfulness. There are countless ways to improve the efficiency of our computing and cloud facilities so long as one simply cares about them - the climate crisis is a modern, global elephant in the room and it should equally be an elephant in project planning. Somebody who does not care about the climate (at their own peril) could create scientific software, could hire servers in a datacentre, could perhaps even build their own datacentre, not considering these factors. But you, the reader, are more careful than that; we should be acutely aware of our electricity-based carbon emissions in the same way we are aware of the plastic we are putting in our general waste or fumes coming out of our cars. It shouldn't be an 'out of sight, out of mind' problem, but we know it is an inevitability that we will produce this waste as part of our lives. Thus, it shouldn't be all or nothing; we should make a careful assessment of where we can reduce our energy use, and where we cannot. We should let it sway our judgements and play a factor in our science and scientific software. That way, it is possible to use the technology of the 'information age' in a way that is beneficial for both people and the planet.

## References

- [1] Piers Forster, Venkatachalam Ramaswamy, Paulo Artaxo, Terje Berntsen, Richard Betts, David W Fahey, James Haywood, Judith Lean, David C Lowe, Gunnar Myhre, et al. Changes in atmospheric constituents and in radiative forcing. chapter 2. In *Climate change 2007. The physical science basis*. Cambridge University Press, 2007.
- [2] Loïc Lannelongue, Jason Grealey, Alex Bateman, and Michael Inouye. Ten simple rules to make your computing more environmentally sustainable, 2021.
- [3] Joshua Aslan, Kieren Mayers, Jonathan G Koomey, and Chris France. Electricity intensity of internet data transmission: Untangling the estimates. *Journal of Industrial Ecology*, 22(4):785–798, 2018.
- [4] iphone 11 product environmental report, 2019. URL [https://www.apple.com/environment/pdf/products/iphone/iPhone\\_11\\_PER\\_sept2019.pdf](https://www.apple.com/environment/pdf/products/iphone/iPhone_11_PER_sept2019.pdf). Accessed: 2022-02-09.

- [5] Greenhouse gas reporting: conversion factors 2020, 2020. URL <https://www.gov.uk/government/publications/greenhouse-gas-reporting-conversion-factors-2020>. Accessed: 2022-02-09.
- [6] Arvind Narayanan, Joseph Bonneau, Edward Felten, Andrew Miller, and Steven Goldfeder. Bitcoin and cryptocurrency technologies: A comprehensive introduction, 2017.
- [7] Jamie Redman. Nft economy grows exponential: \$1m in non-fungible token sales last week, 2020. URL <https://news.bitcoin.com/nft-economy-grows-exponential-1m-in-non-fungible-token-sales-last-week/>. Accessed: 2021-12-23.
- [8] Memo Atken. The unreasonable ecological cost of #cryptoart, 2020. URL <https://memoakten.medium.com/the-unreasonable-ecological-cost-of-cryptoart-2221d3eb2053>. Accessed: 2021-12-23.
- [9] Jiahui Qiu. Nft’s environmental impact, 2021. URL <https://earth.org/nfts-environmental-impact/>. Accessed: 2021-12-23.
- [10] Fahad Saleh. Blockchain without waste: Proof-of-stake. *The Review of Financial Studies*, 34(3):1156–1190, 2021.
- [11] Loïc Lannelongue, Jason Grealey, and Michael Inouye. Green algorithms: Quantifying the carbon footprint of computation. *Advanced Science*, page 2100707, 2021.
- [12] Jason G Grealey, Loïc Lannelongue, Woeiyuh Saw, Jonathan Marten, Guillaume Méric, Sergio Ruiz Carmona, and Michael Inouye. The carbon footprint of bioinformatics. *bioRxiv*, 2021.
- [13] Niklas Heer. Speed comparison of various programming languages, 2018. URL <https://github.com/niklas-heer/speed-comparison>. Accessed: 2021-12-28.
- [14] Arman Shehabi, Sarah Smith, Dale Sartor, Richard Brown, Magnus Herrlin, Jonathan Koomey, Eric Masanet, Nathaniel Horner, Inês Azevedo, and William Lintner. United states data center energy usage report. 2016.
- [15] Liz Waters. Energy consumption in the uk (ecuk) 1970 to 2019. *ECUK*, 2019.
- [16] Dell powerstore9000t pcf datasheet, 2021. URL <https://www.delltechnologies.com/asset/en-us/products/storage/technical-support/powerstore9000t-pcf-datasheet.pdf>. Accessed: 2022-02-04.
- [17] Richard L Moore, Jim D’Aoust, Robert H McDonald, and David Minor. Disk and tape storage cost models. In *Archiving Conference*, volume 1, pages 29–32. Society for Imaging Science and Technology, 2007.
- [18] Fergus Imrie, Anthony R Bradley, and Charlotte M Deane. Generating property-matched decoy molecules using deep learning. *Bioinformatics*, 37(15):2134–2141, 2021.
- [19] Kaushik Bhattacharya, Bamdad Hosseini, Nikola B Kovachki, and Andrew M Stuart. Model reduction and neural networks for parametric pdes. *arXiv preprint arXiv:2005.03180*, 2020.
- [20] Li Deng. The mnist database of handwritten digit images for machine learning research. *IEEE Signal Processing Magazine*, 29(6):141–142, 2012.
- [21] Bryan Klimt and Yiming Yang. The enron corpus: A new dataset for email classification research. In *European Conference on Machine Learning*, pages 217–226. Springer, 2004.
- [22] Emma Strubell, Ananya Ganesh, and Andrew McCallum. Energy and policy considerations for deep learning in nlp. *arXiv preprint arXiv:1906.02243*, 2019.
- [23] David Findlay. *A Practical Guide to the ISIS Neutron and Muon Source*. UKRI Creative Services, 2021.



- [24] Eric Masanet, Arman Shehabi, Lavanya Ramakrishnan, Jiaqi Liang, Xiaohui Ma, Benjamin Walker, Valerie Hendrix, and Pradeep Mantha. The energy efficiency potential of cloud-based software: A u.s. case study. 2013.
- [25] Eric Masanet, Arman Shehabi, Nuo Lei, Sarah Smith, and Jonathan Koomey. Recalibrating global data center energy-use estimates. *Science*, 367(6481):984–986, 2020.
- [26] Urs Hölzle. Data centers are more energy efficient than ever. 2020. URL <https://www.blog.google/outreach-initiatives/sustainability/data-centers-energy-efficient>.
- [27] Christina Ospina. Cooling your home but warming the planet: How we can stop air conditioning from worsening climate change. *Climate Institute. New York*, page 9, 2018.
- [28] Richard Evans and Jim Gao. Deepmind ai reduces google data centre cooling bill by 40 URL <https://deepmind.com/blog/article/deepmind-ai-reduces-google-data-centre-cooling-bill-40/>. Accessed: 2022-02-09.
- [29] Andika Candra Jaya, Cutifa Safitri, and Rila Mandala. Sneakernet: A technological overview and improvement. In *2020 IEEE International Conference on Sustainable Engineering and Creative Computing (ICSECC)*, pages 287–291. IEEE, 2020.
- [30] Mercedes-benz sprinter promotional brochure, 2019. URL <https://www.mercedes-benz.co.uk/vans/content/dam/vans/united-kingdom/sprinter-panel-van/Sprinter%20brochure.pdf>. Accessed: 2022-02-04.
- [31] CJ Hubbard. New british-built arrival electric van: in-depth info, official pics and more, 2021. URL <https://www.parkers.co.uk/vans-pickups/news/2021/arrival-electric-van/>. Accessed: 2022-02-04.
- [32] Hashem Akbari. Shade trees reduce building energy use and co2 emissions from power plants. *Environmental pollution*, 116:S119–S126, 2002.
- [33] Alia al Ghussain. The biggest problem with carbon offsetting is that it doesn’t really work. *Greenpeace*, 2021.
- [34] Bichlien Nguyen, Julie Sinistore, Jake Smith, Praneet S. Arshi, Lauren M. Johnson, Tim Kidman, T.J. diCaprio, Doug Carmean, and Karin Strauss. Architecting datacenters for sustainability: Greener data storage using synthetic dna. In *Electronics Goes Green 2020*. Fraunhofer IZM, IEEE, September 2020. URL <https://www.microsoft.com/en-us/research/publication/architecting-datacenters-for-sustainability-greener-data-storage-using-synthetic-dna/>.